

# Spring Kafka beyond the basics

Lessons learned on our Kafka journey at ING Bank

SPRING I/O BARCELONA – MAY 27<sup>TH</sup> 2022

Tim van Baarsen



# Who am I?



**Tim van Baarsen**  
Software Engineer @ ING Netherlands



The Netherlands

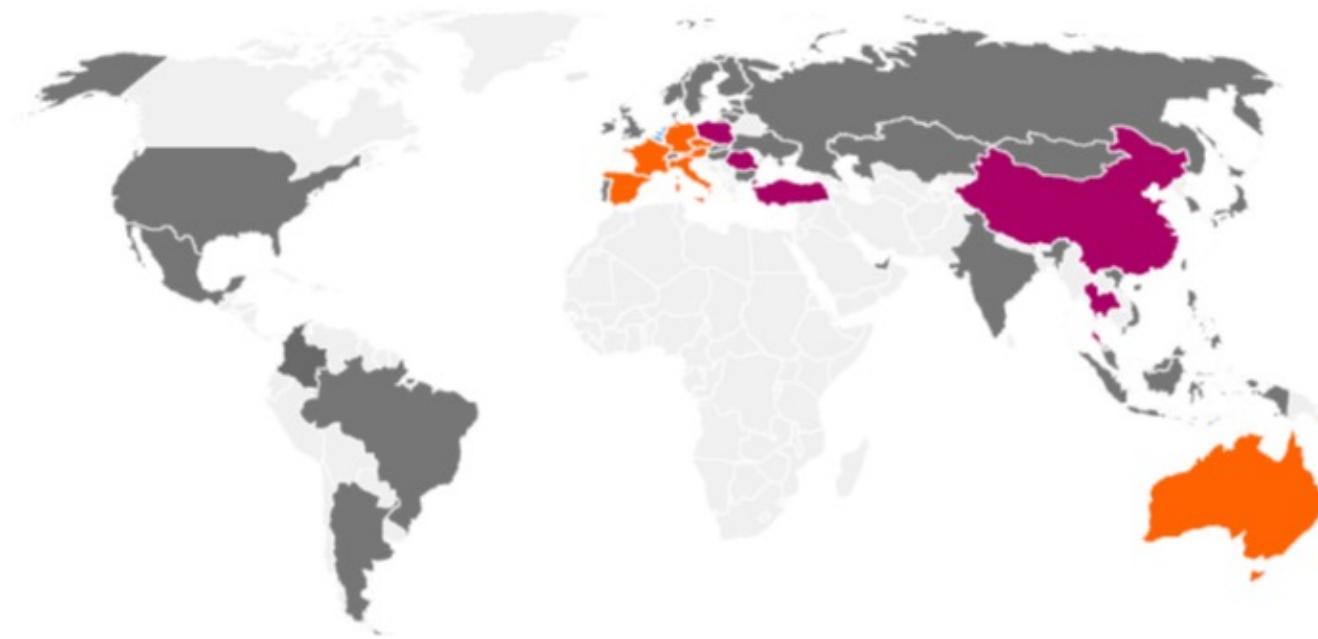


Amsterdam



Team Dora

# ING is active in more than 40 countries



- Market Leaders Benelux
- Growth Markets
- Challengers
- Wholesale Banking

**Disclaimer:** Please note that ING Bank does not have a banking license in the US and is therefore not permitted to conduct banking activities in the US.

# Kafka @ ING

## Frontrunners in Kafka

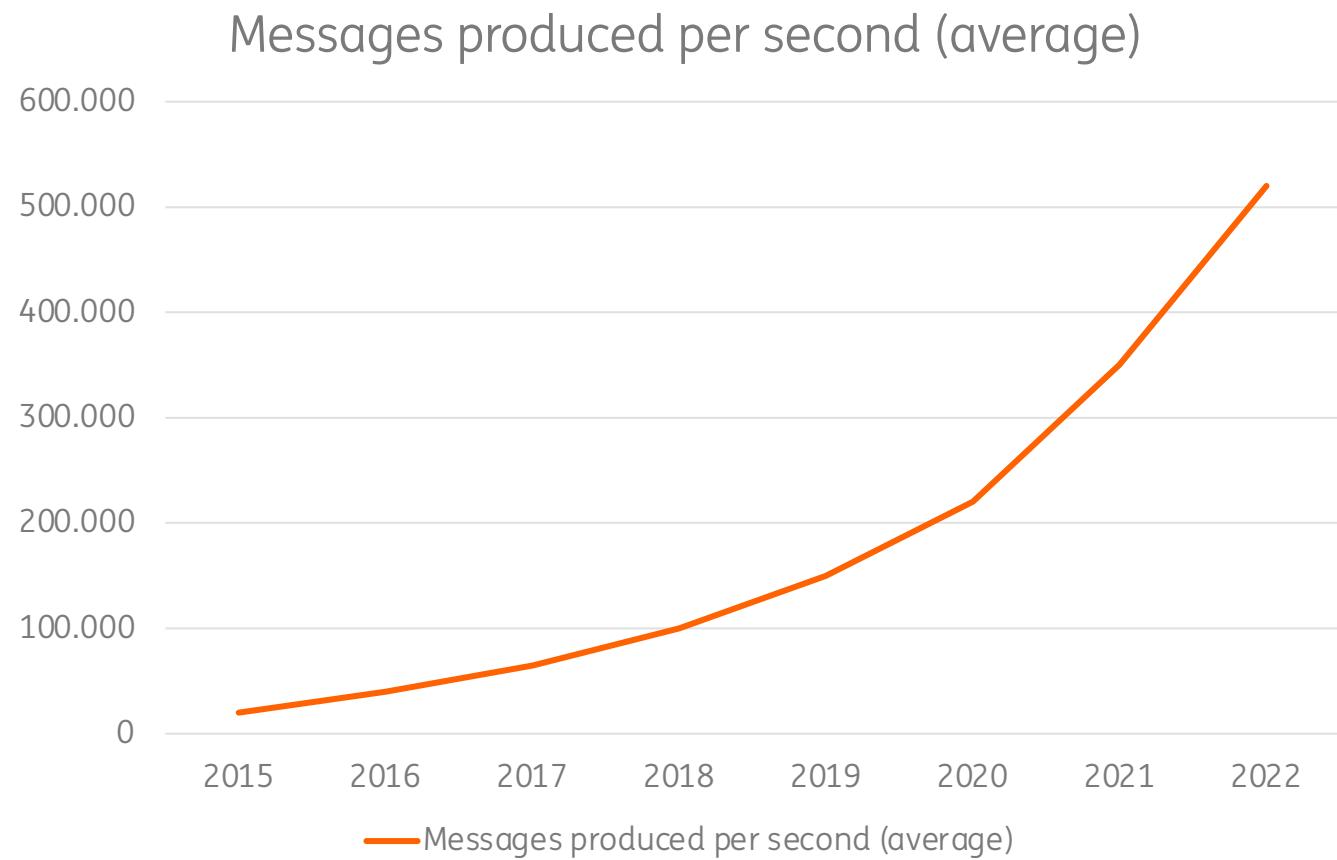
### Running in production:

- 7 years
- 5000+ **topics**
- Serving 850+ Development **teams**
- Self service topic management

The screenshot shows the Touchpoint API & Stream Marketplace interface. At the top, there are tabs for Touchpoint, Console, and a search bar. Below the search bar are navigation links: Discover, Design, Develop, Distribute, and Help. A user profile icon is in the top right. The main area is titled "API & Stream Marketplace" and "Topics". It features two cards: one for "test-topic" (DEV, TST, Logging) and one for "stock-quotes" (DEV, TST, ACC, PRD, Business). A "View all" link is on the right. At the bottom, there's a footer with links to About Touchpoint, Become a partner, and Terms & Conditions, along with a copyright notice: ©2022 - ING.

# Kafka @ ING

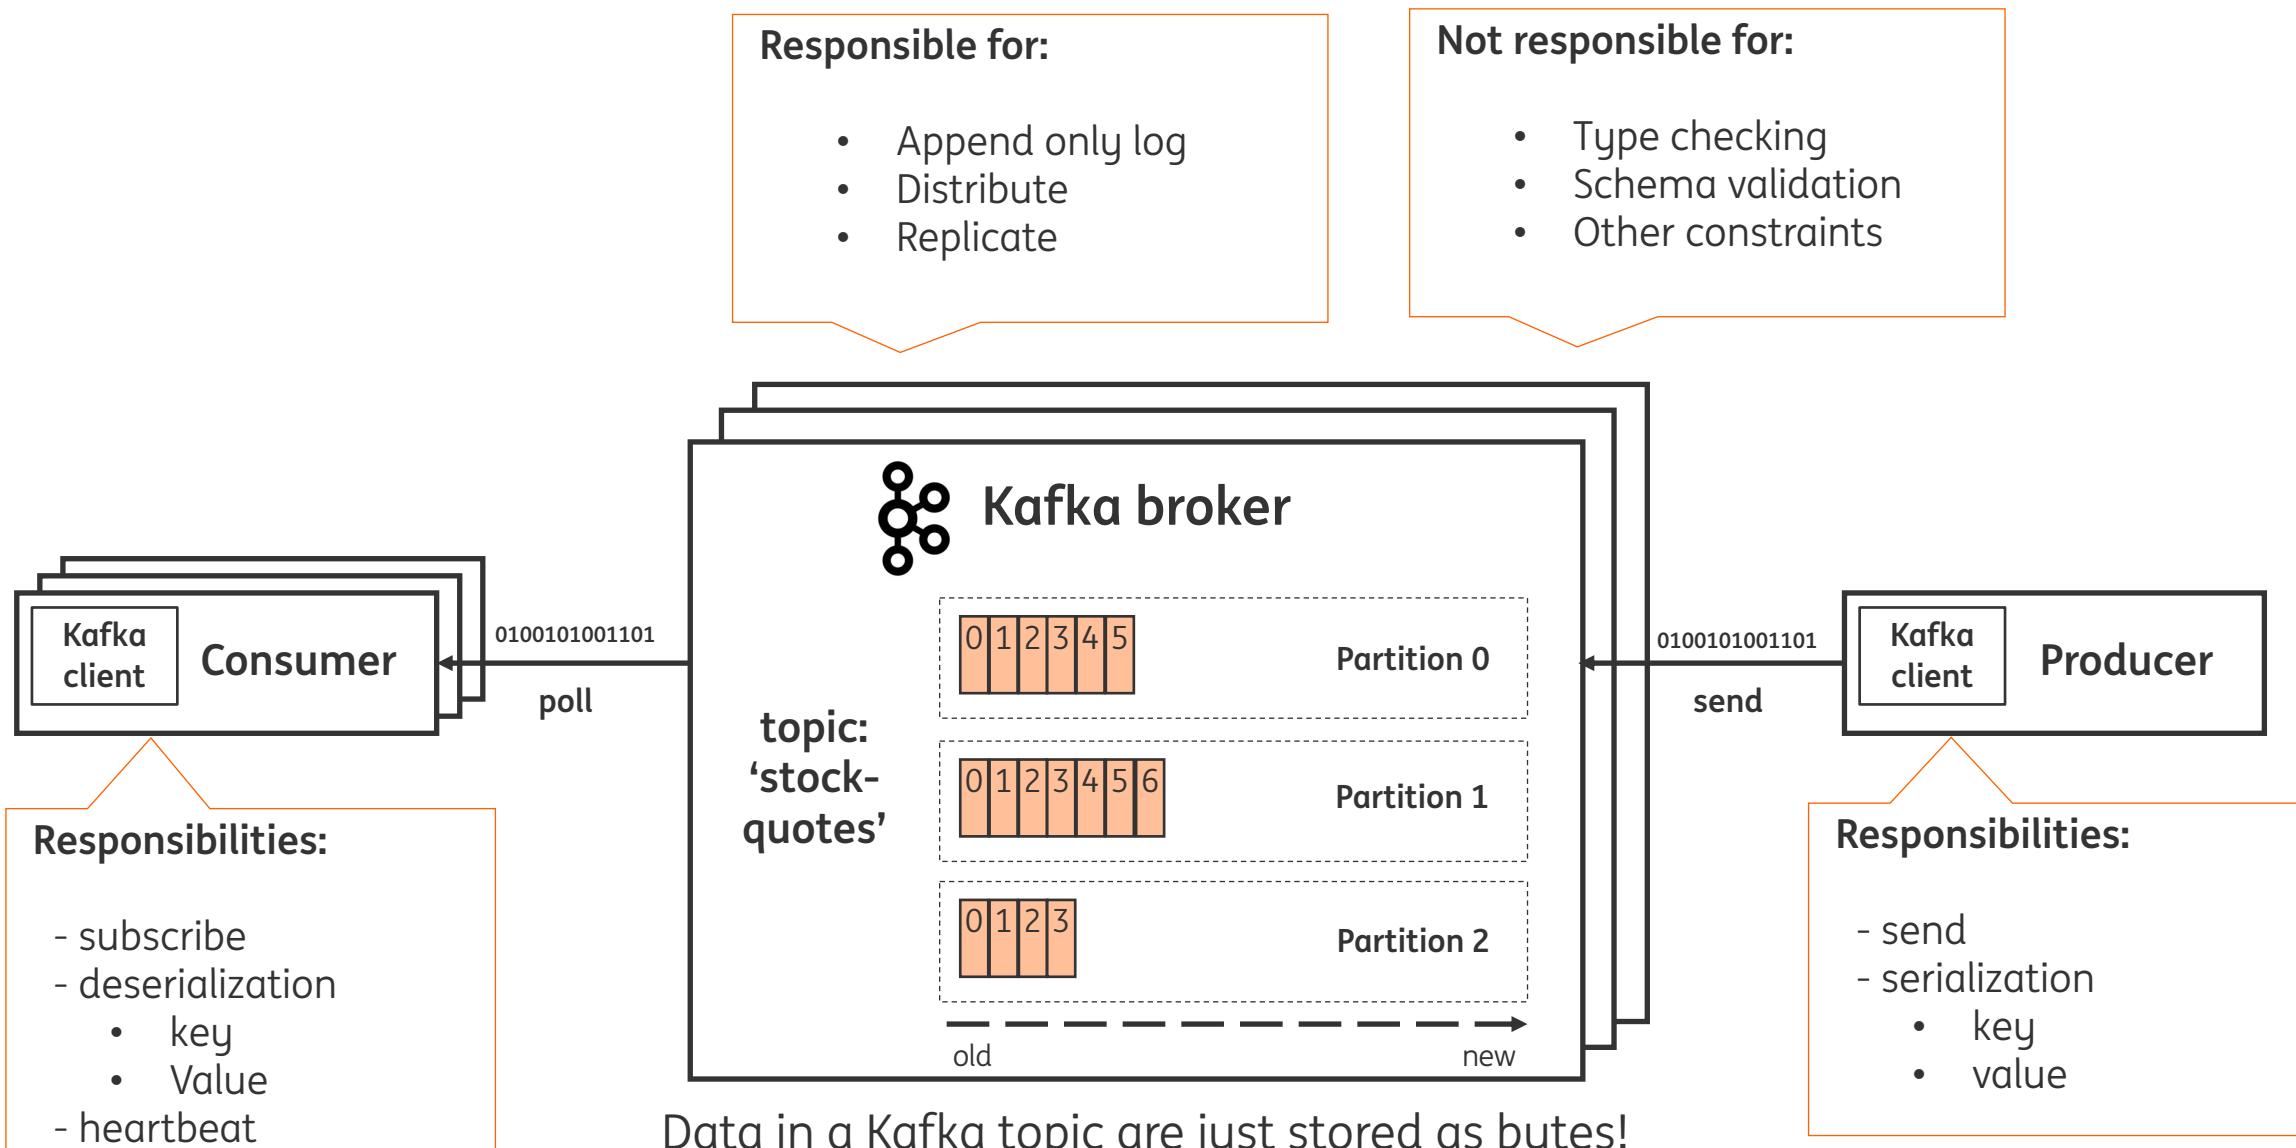
Traffic is growing with 10%+ monthly



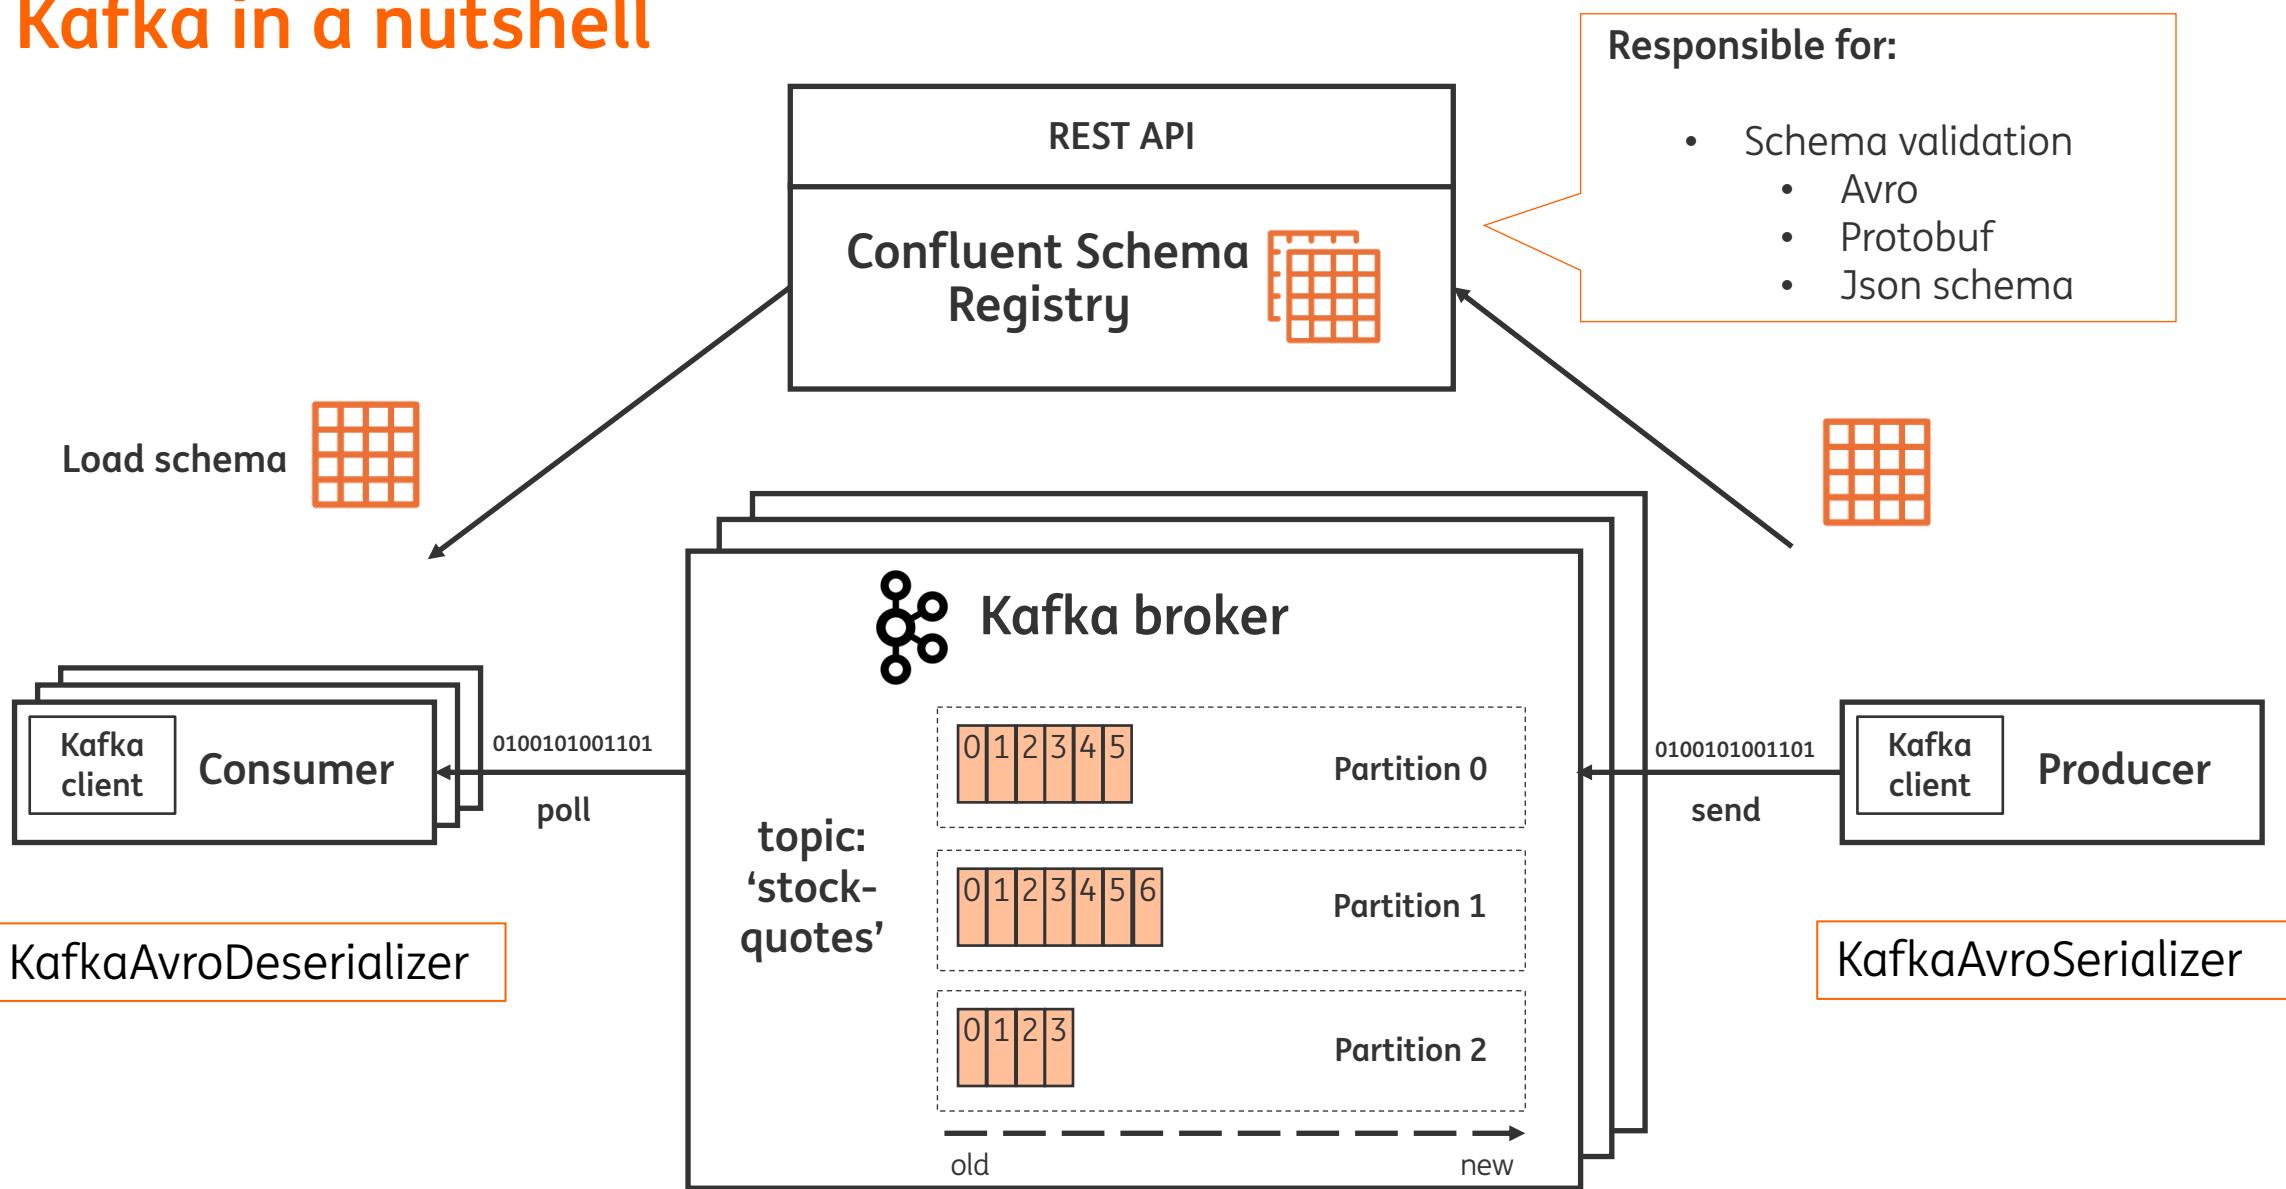
# Agenda

- Kafka in a nutshell
- Spring Kafka essentials
- Scenario: Poison pill / Deserialization exceptions
- Scenario: Lack of exception handling
- Testing
- Monitoring
- Wrap-up
- Questions

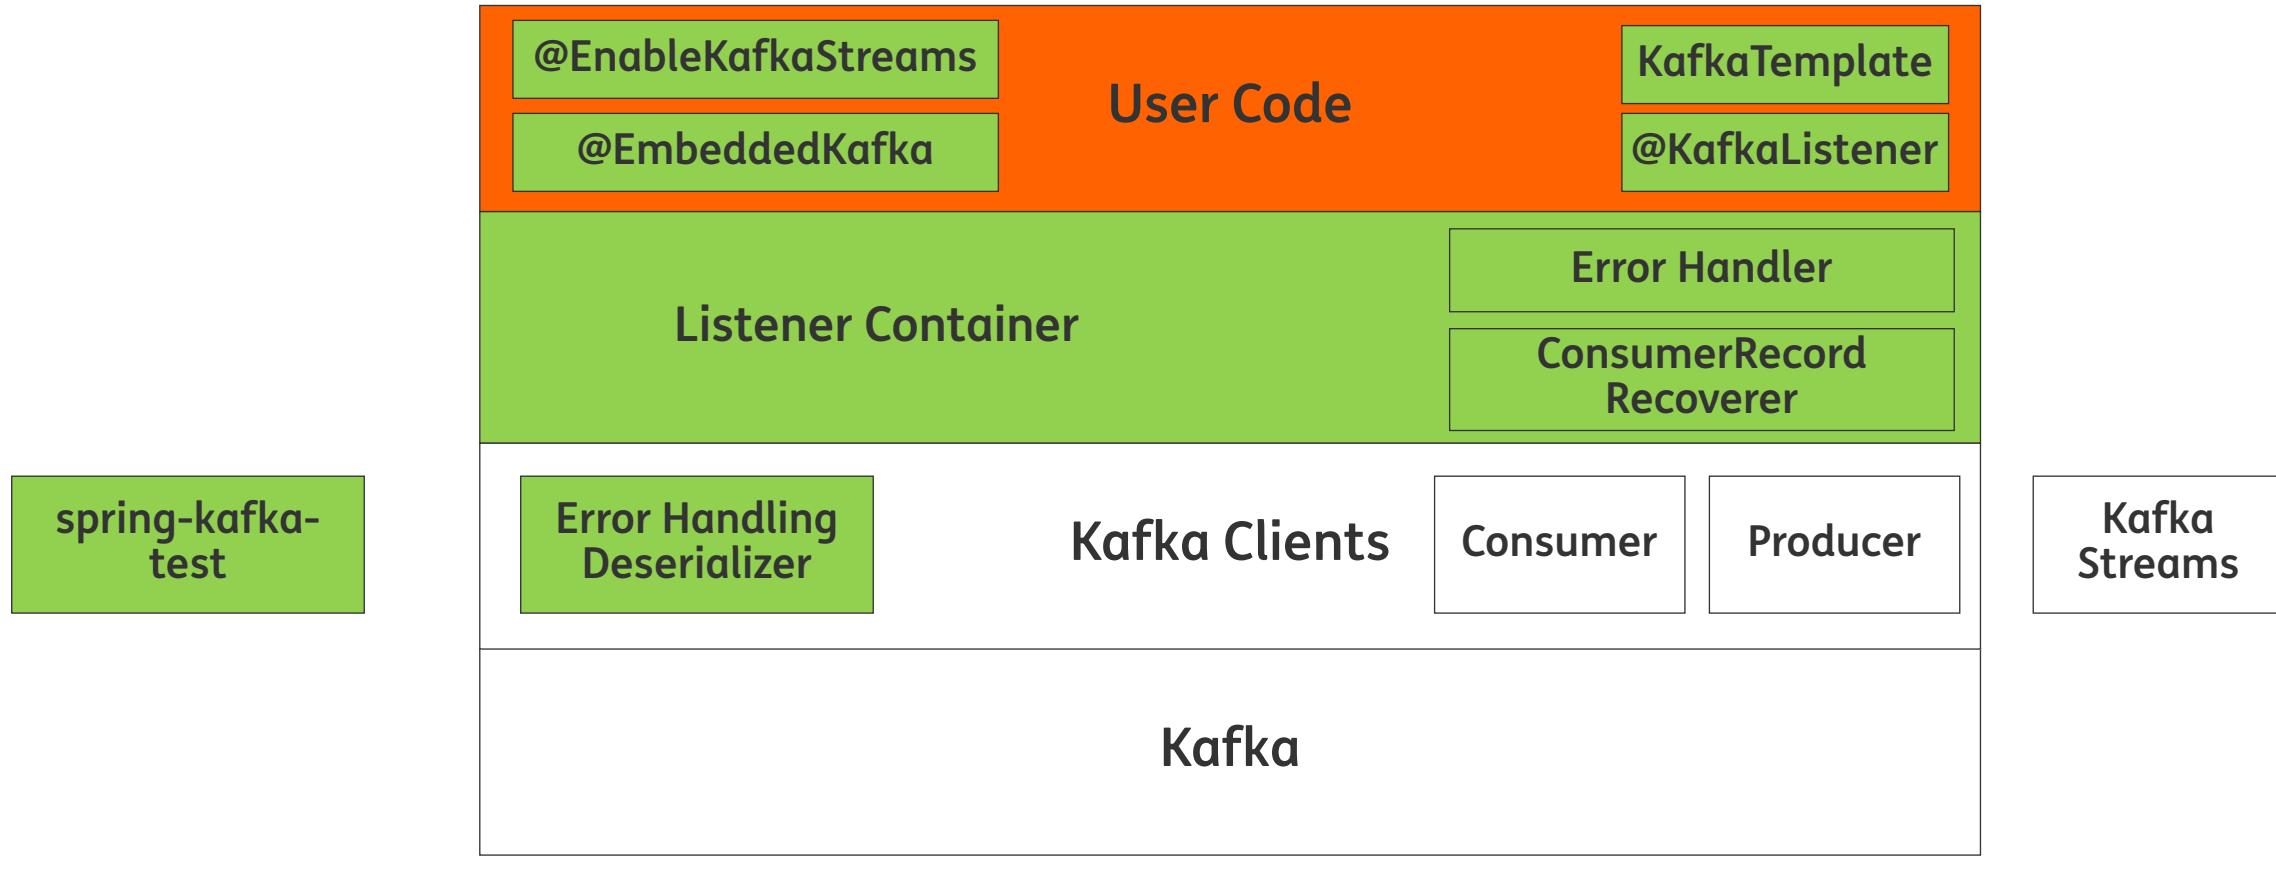
# Kafka in a nutshell



# Kafka in a nutshell



# Spring for Apache Kafka - essentials



# Spring Kafka & Spring Boot - Producer

```
@Component
@Slf4j
public class StockQuoteProducer {

    @Autowired
    private KafkaTemplate<String, StockQuote> kafkaTemplate;

    public void produce(StockQuote stockQuote) {
        kafkaTemplate.send("stock-quotes", stockQuote.getSymbol(), stockQuote);
        log.info("Produced stock quote: {}", stockQuote);
    }
}
```

# Spring Kafka & Spring Boot - Consumer

```
@Component  
@Slf4j  
public class StockQuoteConsumer {  
  
    @KafkaListener(topics = "stock-quotes")  
    public void on(StockQuote stockQuote,  
                  @Header(KafkaHeaders RECEIVED_PARTITION_ID) String partition) {  
        log.info("Consumed from partition: {} value: {}", partition, stockQuote);  
    }  
}
```

# Spring Kafka & Spring Boot – Kafka streams

```
● ● ●  
 @Configuration  
 @EnableKafkaStreams  
 public class KafkaStreamsConfig {  
  
     @Bean  
     public KStream<String, StockQuote> kStream(StreamsBuilder streamsBuilder) {  
  
         KStream<String, StockQuote> branchedStream = new KafkaStreamBrancher<String, StockQuote>()  
             .branch((key, value) -> value.getExchange().equalsIgnoreCase("NYSE"), kStream -> kStream.to("stock-quotes-nyse"))  
             .branch((key, value) -> value.getExchange().equalsIgnoreCase("NASDAQ"), kStream -> kStream.to("stock-quotes-nasdaq"))  
             .branch((key, value) -> value.getExchange().equalsIgnoreCase("AMS"), kStream -> kStream.to("stock-quotes-ams"))  
             .defaultBranch(kStream -> kStream.to("stock-quotes-exchange-other"))  
             .onTopOf(streamsBuilder.stream("stock-quotes"));  
  
         return branchedStream;  
     }  
 }
```

# Spring Kafka & Spring Boot – Configuration (application.yml)

```
spring:
  application:
    name: producer-application

kafka:
  bootstrap-servers: localhost:9092

producer:
  key-serializer: org.apache.kafka.common.serialization.StringSerializer
  value-serializer: io.confluent.kafka.serializers.KafkaAvroSerializer
  client-id: ${spring.application.name}

properties
  schema.registry.url: http://localhost:8081
```

# After the honeymoon phase is over



Traps



Lessons learned  
The hard way!



Tips & Tricks



Apply in your  
own project(s)

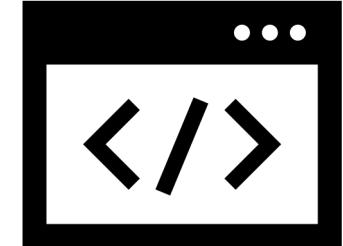
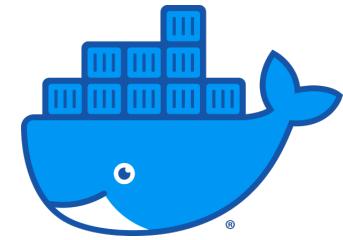


Share with  
your fellow  
developers



# 🚀 Tip: Local development setup

- Docker
  - Kafka Cluster
  - Zookeeper
  - Confluent Schema Registry
- CLI Tools
  - Kafka CLI, Confluent CLI, Kafka cat
- UI Tools
  - Confluent Control Center
  - Kafka UI
  - Conduktor
  - etc





## Scenario: 'Poison Pill' (Deserialization exception)

What is a Poison Pill?



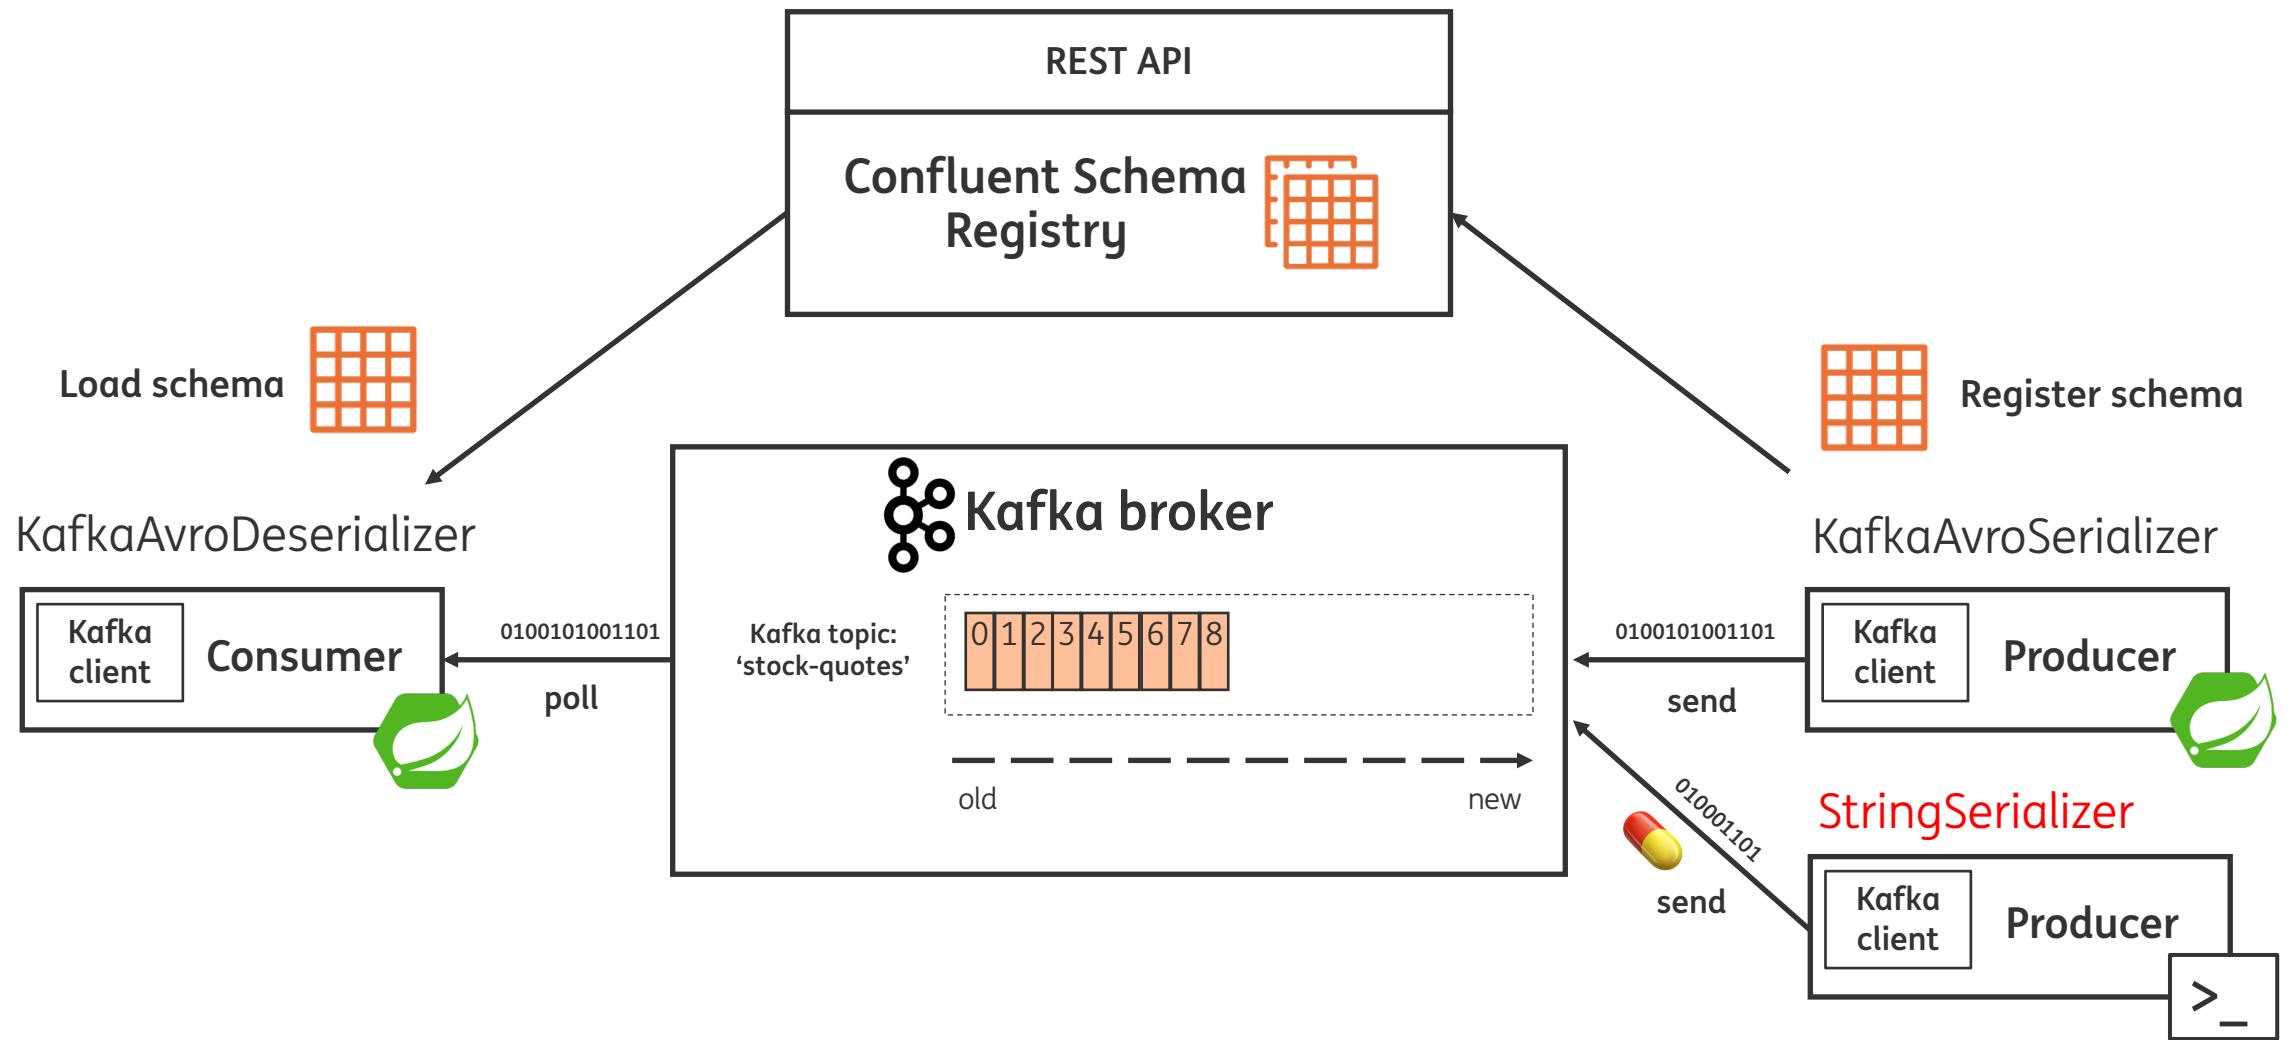
Different forms:

Corrupted  
record

Deserialization  
failure

A **record** that **always fails** when **consumed**,  
no matter how many times it is attempted.

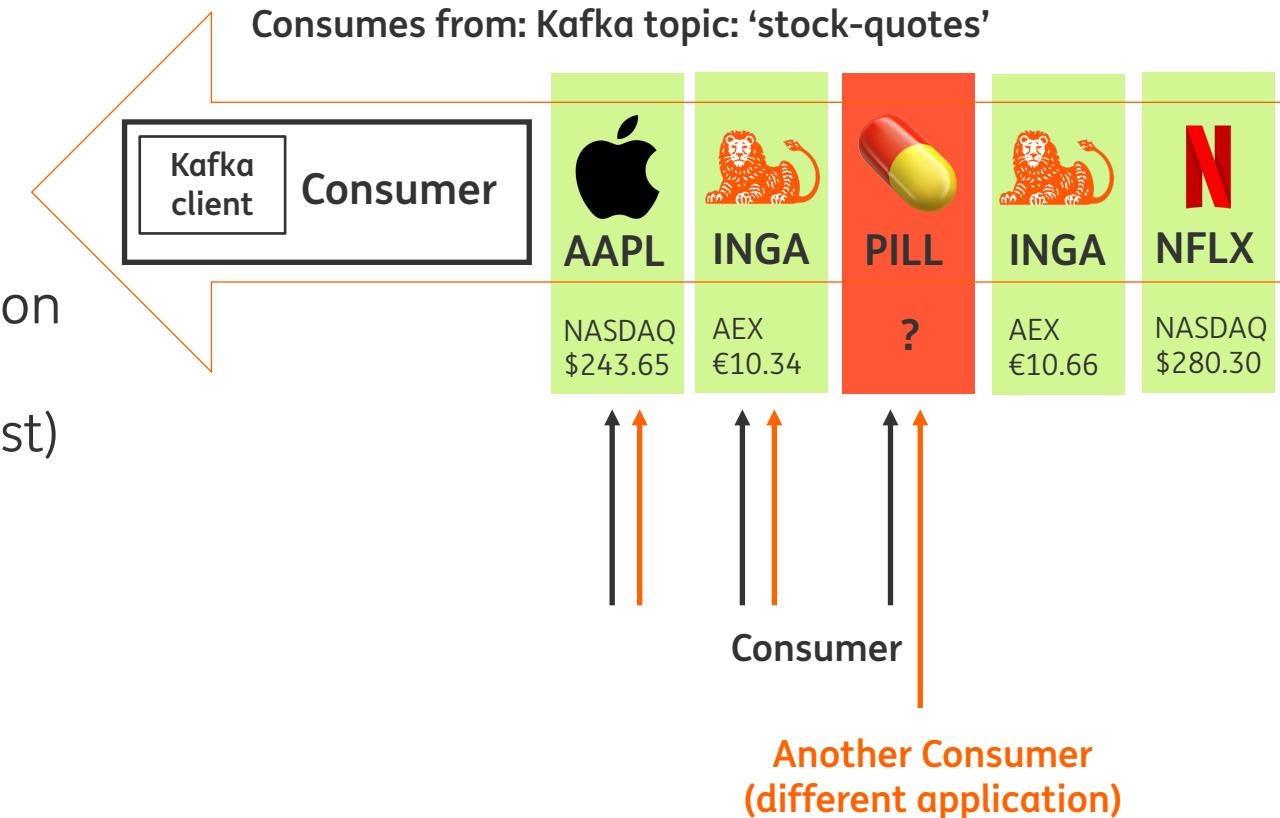
# 💣 Scenario: 'Poison Pill' (Deserialization exception)



# 💣 Scenario: 'Poison Pill' (Deserialization exception)

## Scenario:

- **Consumer** of topics
- Someone **produced** a 'poison' message
  - Consumer fails to deserialize
- **Consequence**
  - Blocks consumption of the topic/partition
  - application can't 'swallow the pill'
  - Try again and again and again (very fast)
  - Log line for every failure



# 💣 Scenario: 'Poison Pill' (Deserialization exception)

- **Result:** log file will grow very fast, flood your disc
- **Impact:** High

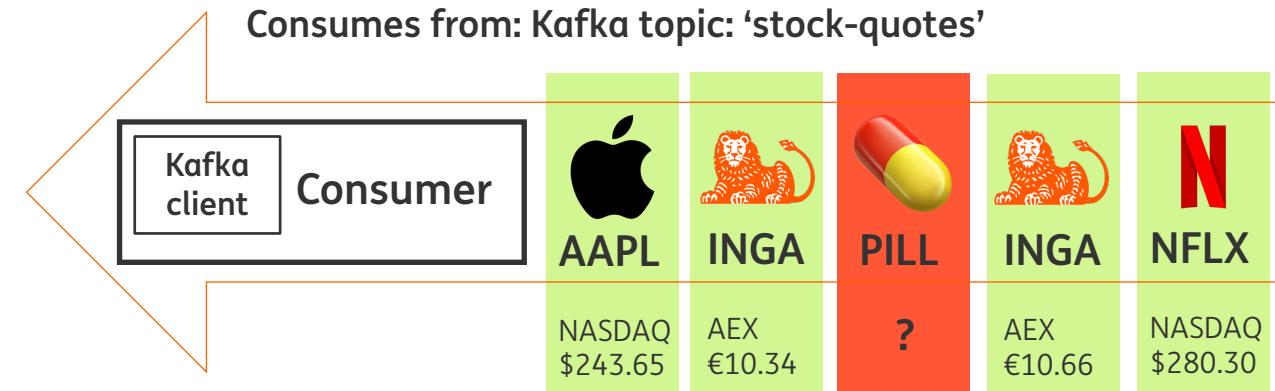
- **How to survive this scenario?**

👎 • Wait until the retention period of the topic has passed

👎 • Change consumer group

👎 • Manually / Programmatically update the offset

👍 • Configure ErrorHandlingDeserializer (Provided by Spring Kafka)



## 💣 Scenario: 'Poison Pill' - Demo





# Scenario: Lack of proper exception handling

## Scenario:

- Consumer
- Exception is thrown in the method handling the message

## **Result (by default)**

- Records that fail are:
  - Retried
  - Logged
  - We move on to the next one



```
 @KafkaListener(topics = "stock-quotes")
 public void on(StockQuote stockQuote) {
     if ("KABOOM".equalsIgnoreCase(stockQuote.getSymbol())) {
         throw new RuntimeException("Whoops something went wrong...");
     }
 }
```

## **Consequence**

- You lose that message!
- Not acceptable in many use-cases!

# 💣 Scenario: Lack of proper exception handling

## Impact

- Dependents on your use-case

## How to survive this scenario?

- Replace / configure DefaultErrorHandler by:
  - CommonLoggingErrorHandler
  - ContainerStoppingErrorHandler
- Use backoff strategy for recoverable exceptions
- Configure ConsumerRecordRecoverer
  - Dead letter topic
  - Implement your own recoverer

## ⚠ Scenario: Lack of proper exception handling - Demo



# Testing

- Support for Integration testing `@EmbeddedKafka`

```
@EmbeddedKafka
@SpringJUnitConfig
public class EmbeddedKafkaIntegrationTest {

    @Autowired
    private EmbeddedKafkaBroker embeddedKafkaBroker;

    @Test
    public void yourTestHere() throws Exception {
        // TODO implement ;
    }
}
```

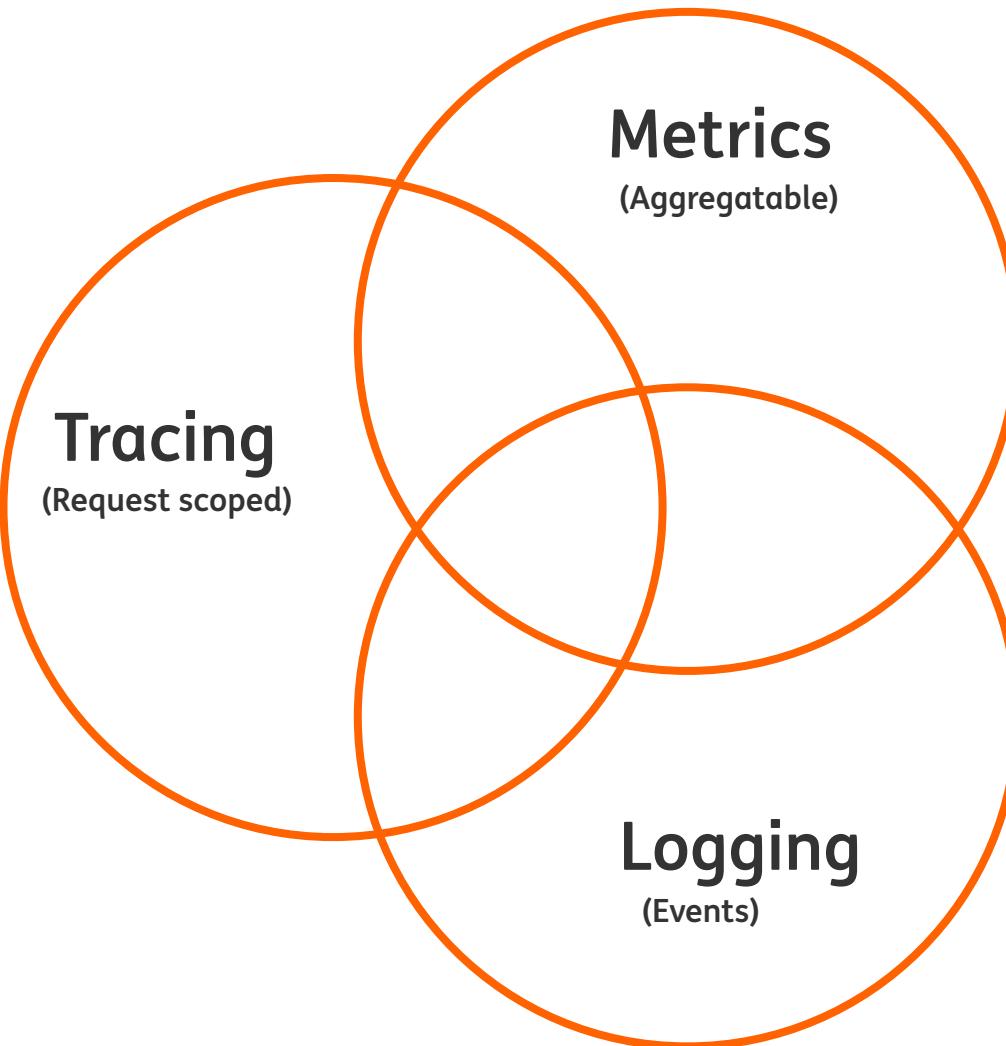
# Testing – Alternatives to EmbeddedKafka

- Test containers
- Focus on unit test first
- Kafka streams
  - Topology test driver

# Testing – Test topology test driver

```
● ● ●  
@Test  
void stockQuoteFromAmsterdamStockExchangeEndUpOnTopicQuotesAmsTopic() {  
    StockQuote stockQuote = new StockQuote("INGA", "AMS", "10.99", "EUR",  
                                         "Description", Instant.now());  
  
    stockQuoteInputTopic.pipeInput(stockQuote.getSymbol(), stockQuote);  
  
    assertThat(stockQuoteAmsOutputTopic.isEmpty()).isFalse();  
    assertThat(stockQuoteAmsOutputTopic.getQueueSize()).isEqualTo(1L);  
    assertThat(stockQuoteAmsOutputTopic.readValue()).isEqualTo(stockQuote);  
  
    assertThat(stockQuoteNyseOutputTopic.isEmpty()).isTrue();  
    assertThat(stockQuoteNasdaqOutputTopic.isEmpty()).isTrue();  
    assertThat(stockQuoteOtherOutputTopic.isEmpty()).isTrue();  
}
```

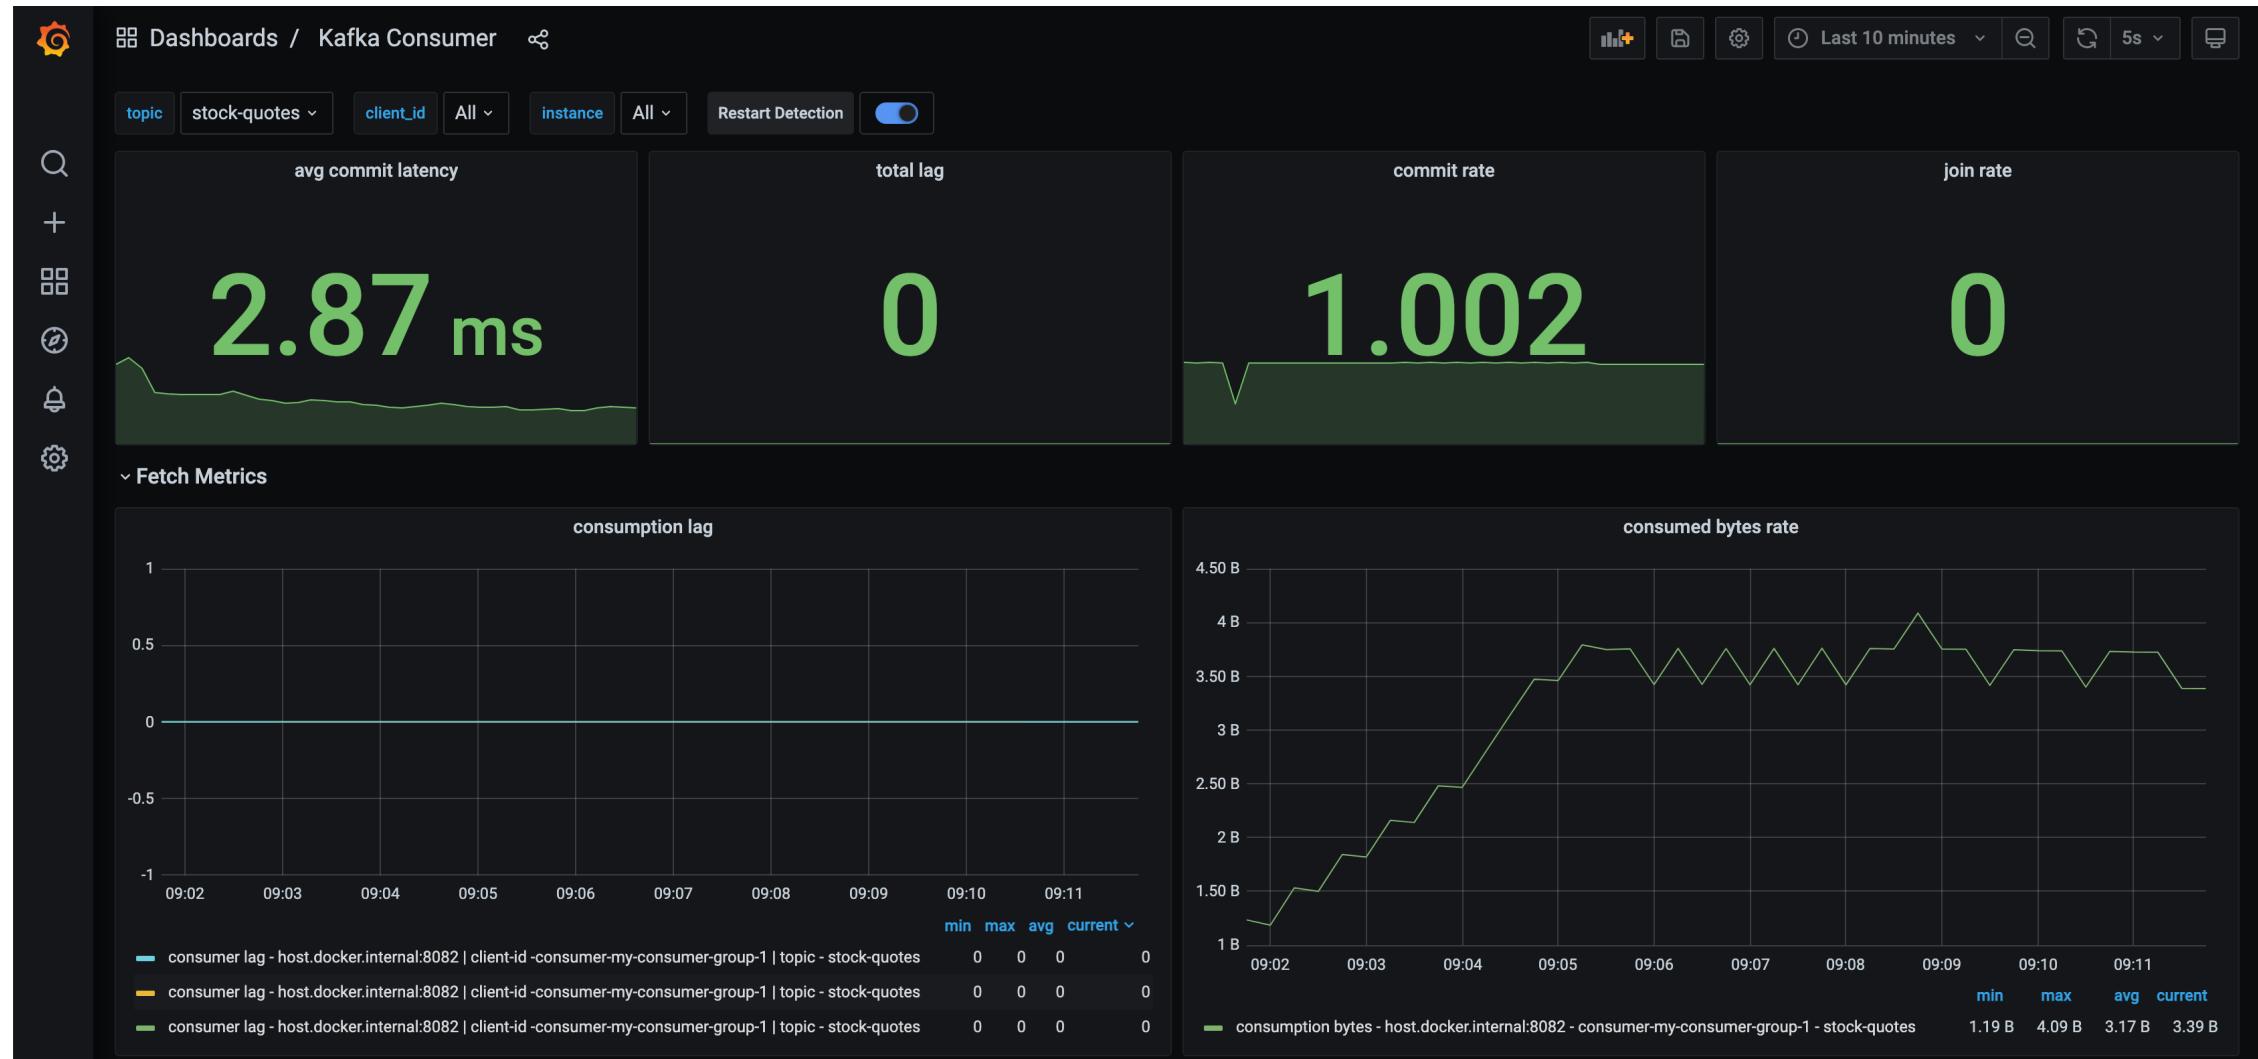
# Monitoring - Three Pillars observability



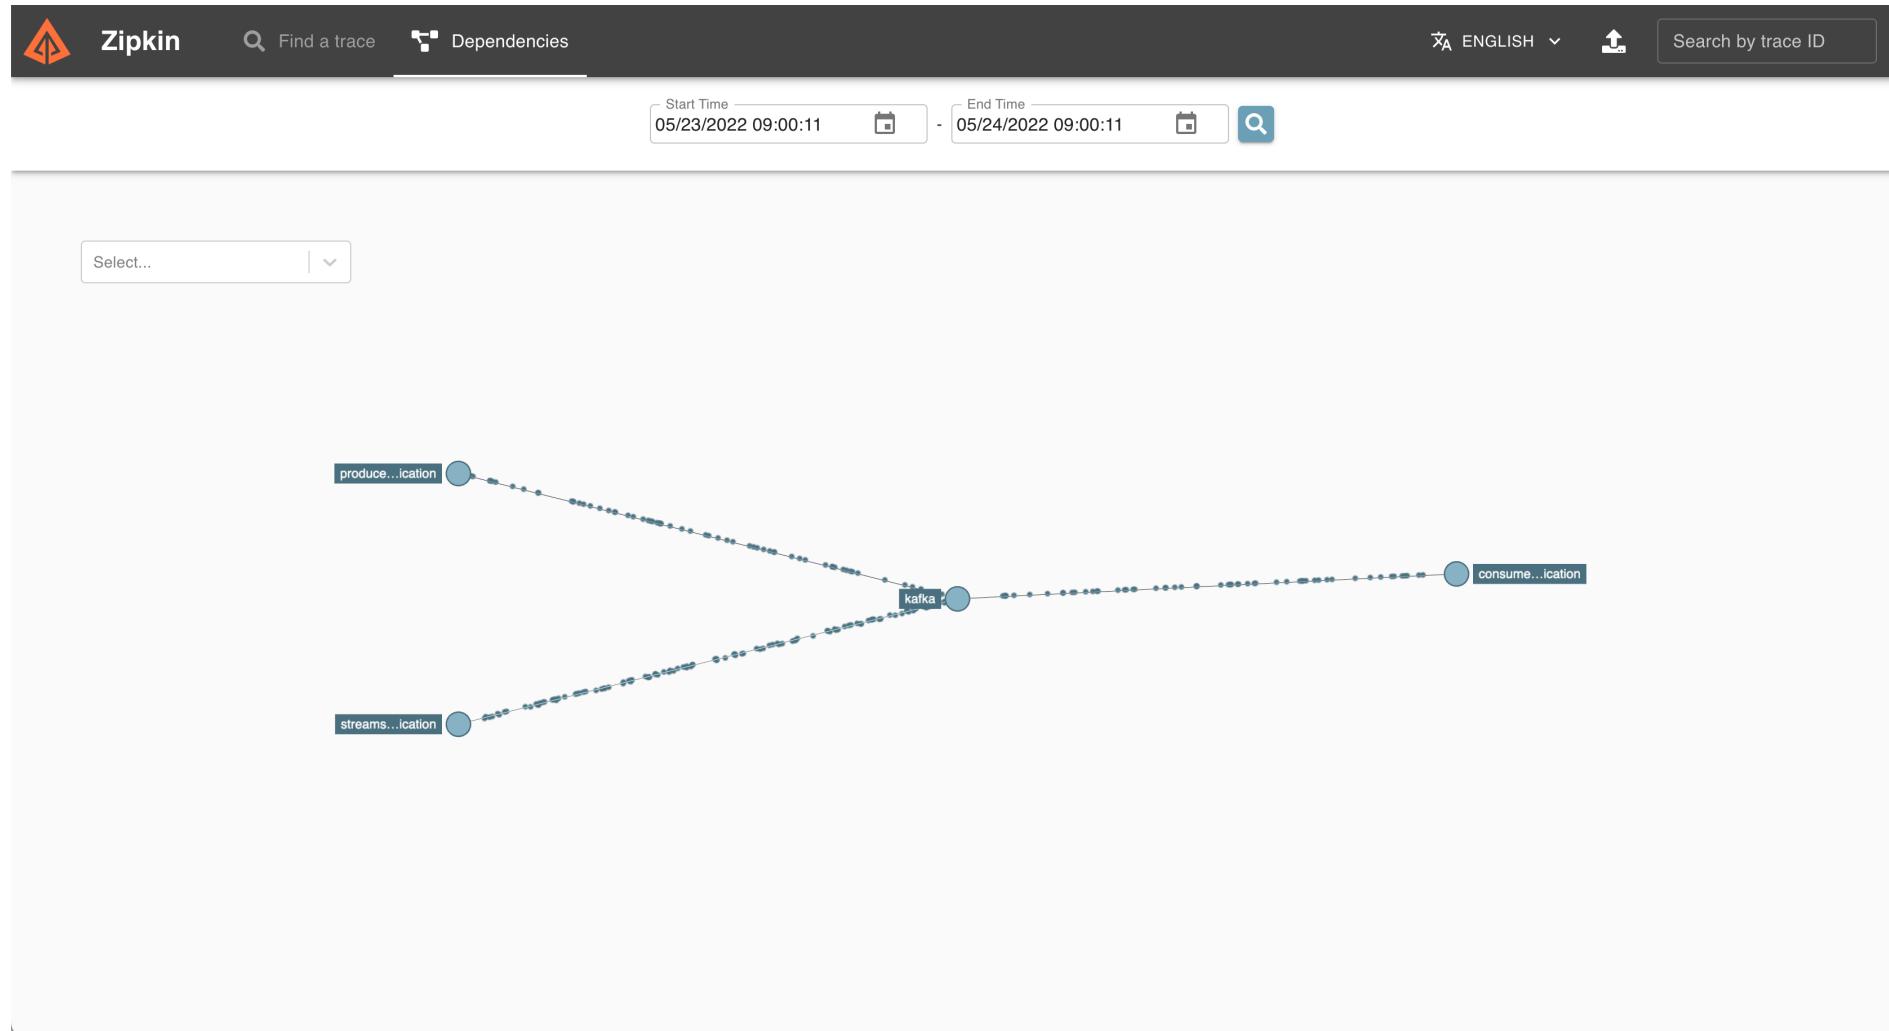
# Monitoring – Metrics

- Out of the box Kafka metrics
    - consumers
    - producer
    - streams
  - Micrometer
  - Spring Boot Actuator Metrics

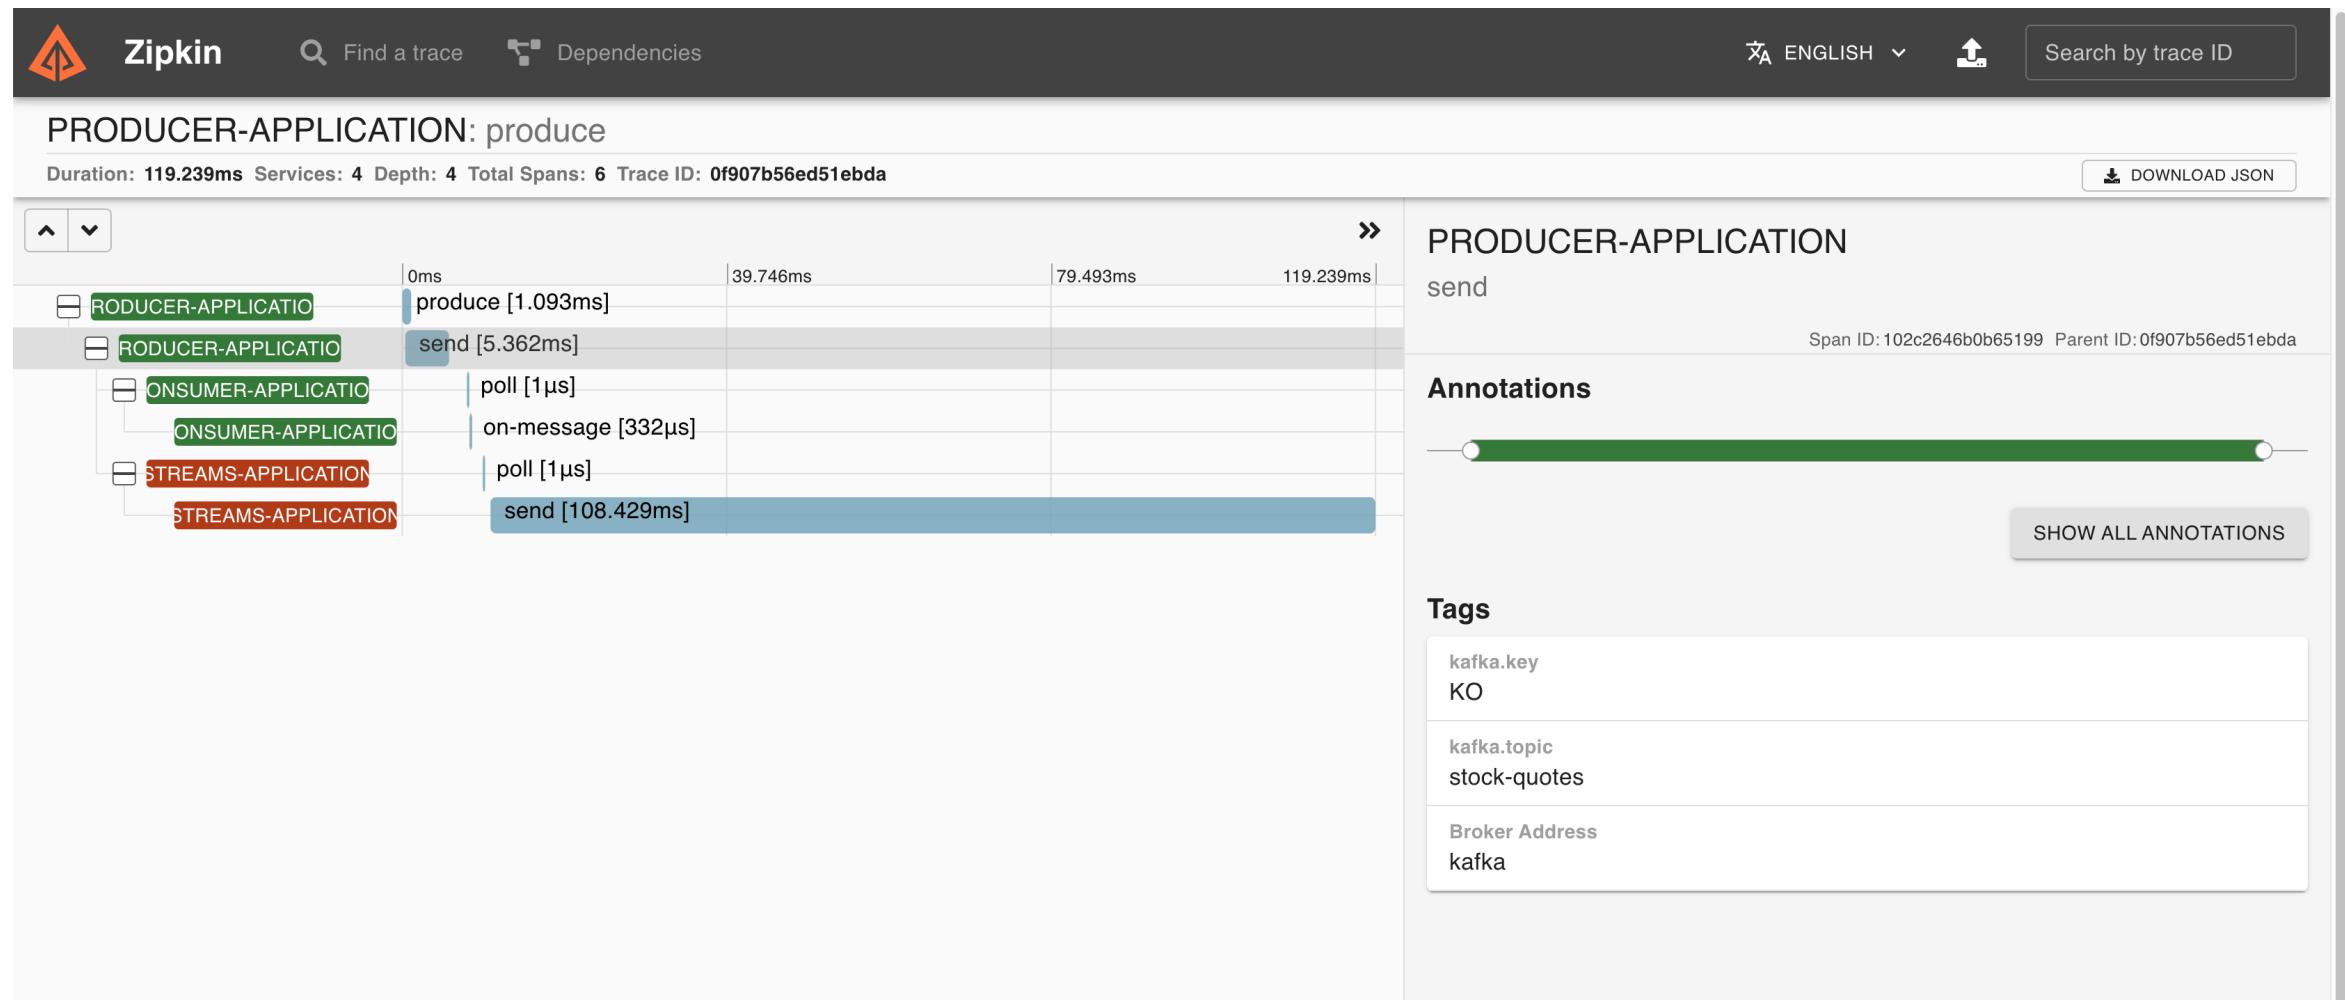
# Monitoring – Metrics



# Monitoring – Distributed tracing



# Monitoring – Distributed tracing





## Lessons learned

- Invest time in your **local development** environment (fast feedback loop)
- Use **Spring Kafka** but also understand the core **Kafka APIs**
- Consumer:
  - Expect the unexpected!
    - Handle Deserialization exceptions a.k.a. 'poison pills'
    - Proper exception handling
    - Validate incoming data
- Producer:
  - Don't change your serializers
  - Leverage Apache Avro + Confluent Schema registry
  - Don't break compatibility for your consumers!



## Lessons learned

- Security
  - Who can produce data to your topics?
  - Who can consume?
- Monitor your topics, consumer groups & applications in production
  - Micrometer
  - Spring Cloud Sleuth
- Don't overdo integration test!

# Questions



# Thanks for joining my talk!



@TimvanBaarsen



<https://www.medium.com/@TimvanBaarsen>



<https://github.com/j-tim/spring-io-barcelona-2022-spring-kafka-beyond-the-basics>